


GSYNC : A GPGPU Synchronization library

CS527 – Parallel Computer Architecture

CSD3171 – Georgios Anagnopoulos



Summary

GSYNC

GSYNC Overview:

- GPU Locks
- GPU Barriers
- Merge library with M4 MACROS

- Summary

- Background

- GSYNC

- Results

- Conclusion

The Pascal GP102 Architecture:

GP102 is composed of :

- Graphics Processing Clusters (GPCs)
- Streaming Multiprocessors (SMs)
- and Memory Controllers.

GPU	GeForce GTX 1080 (Pascal)
SMs	20
CUDA Cores	2560
Base Clock	1607 MHz
GPU Boost Clock	1733 MHz
GFLOPs	8873 ¹
Texture Units	160
Texel fill-rate	277.3 Gigatexels/sec
Memory Clock (Data Rate)	10,000 MHz
Memory Bandwidth	320 GB/sec
ROPs	64
L2 Cache Size	2048 KB
TDP	180 Watts
Transistors	7.2 billion
Die Size	314 mm²
Manufacturing Process	16 nm

Summary
• Background
GSYNC
Results
Conclusion

Background

- Summary
- Background
- GSYNC
- Results
- Conclusion

Pascal Architecture

GP100 Chip

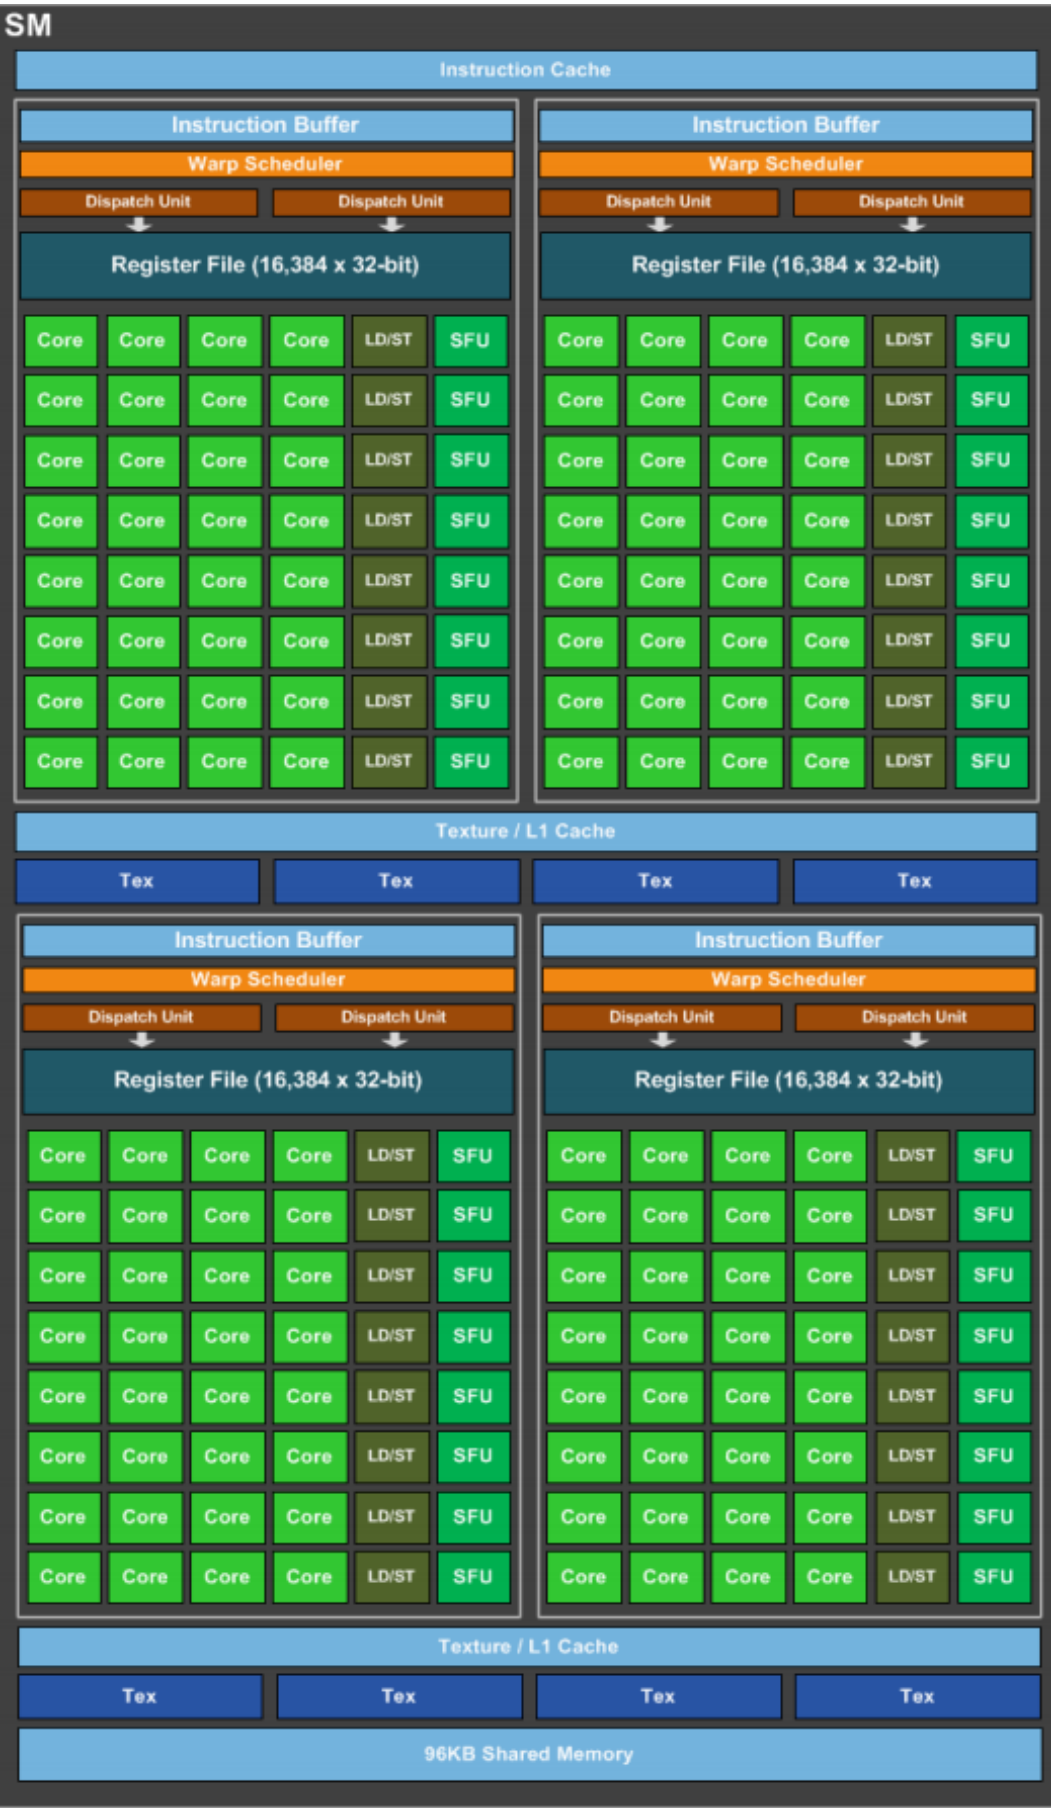


Background

GPU Hardware Architecture



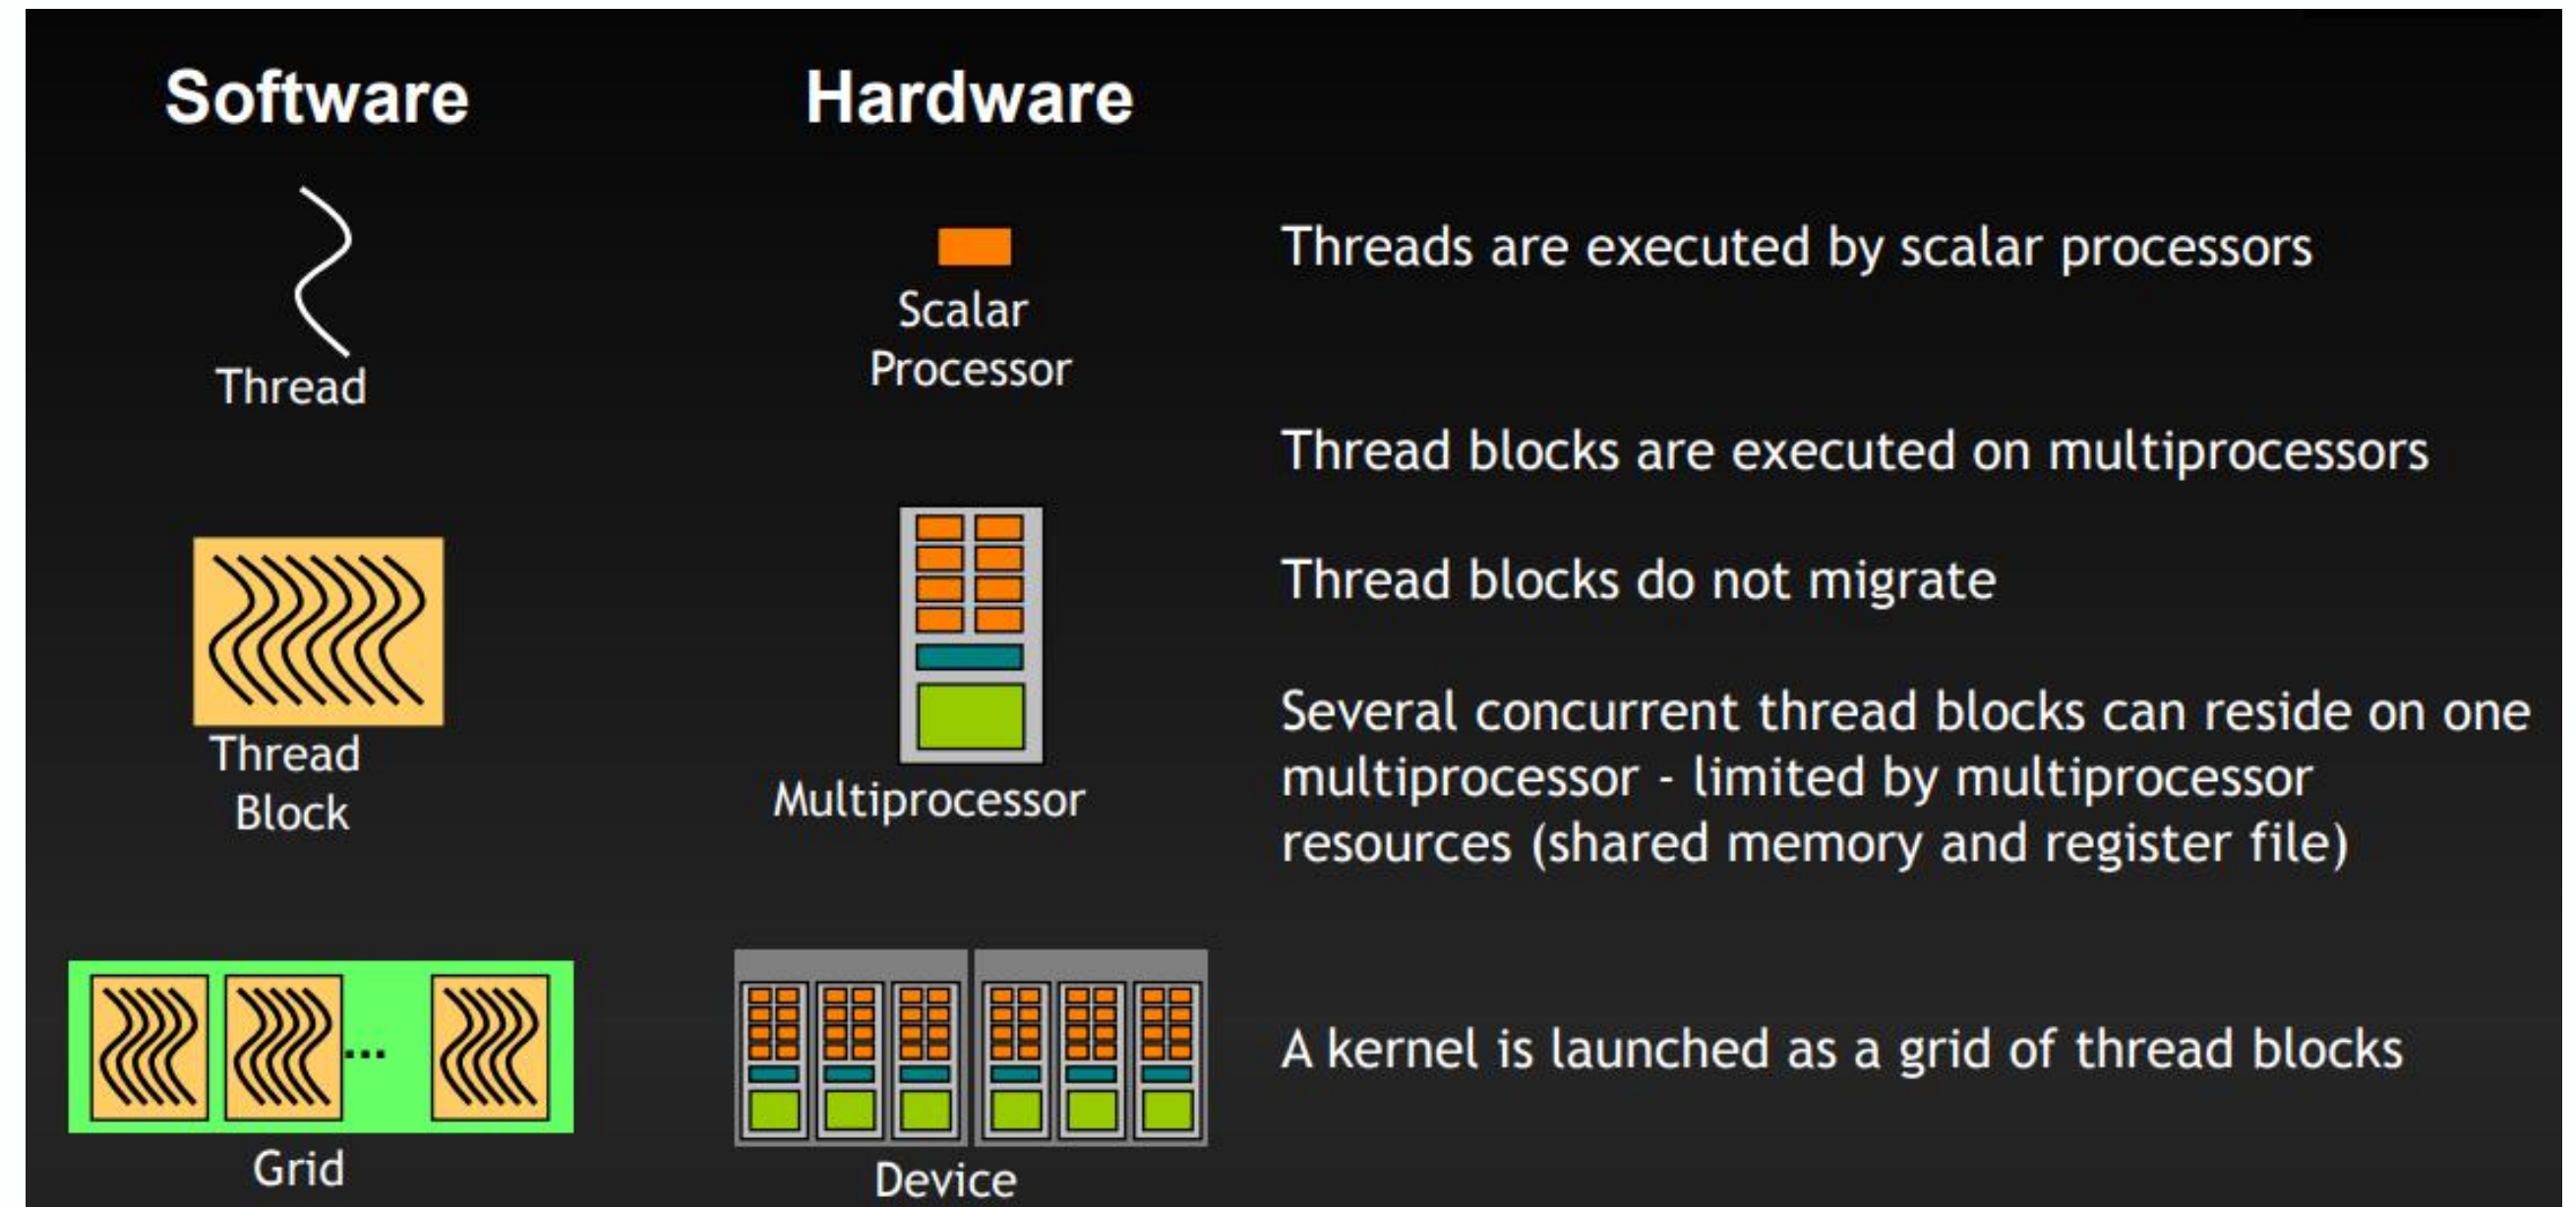
Streaming
Multiprocessor



- Summary
- Background
- GSYNC
- Results
- Conclusion

Background

GPU Hardware Architecture



Summary

• Background

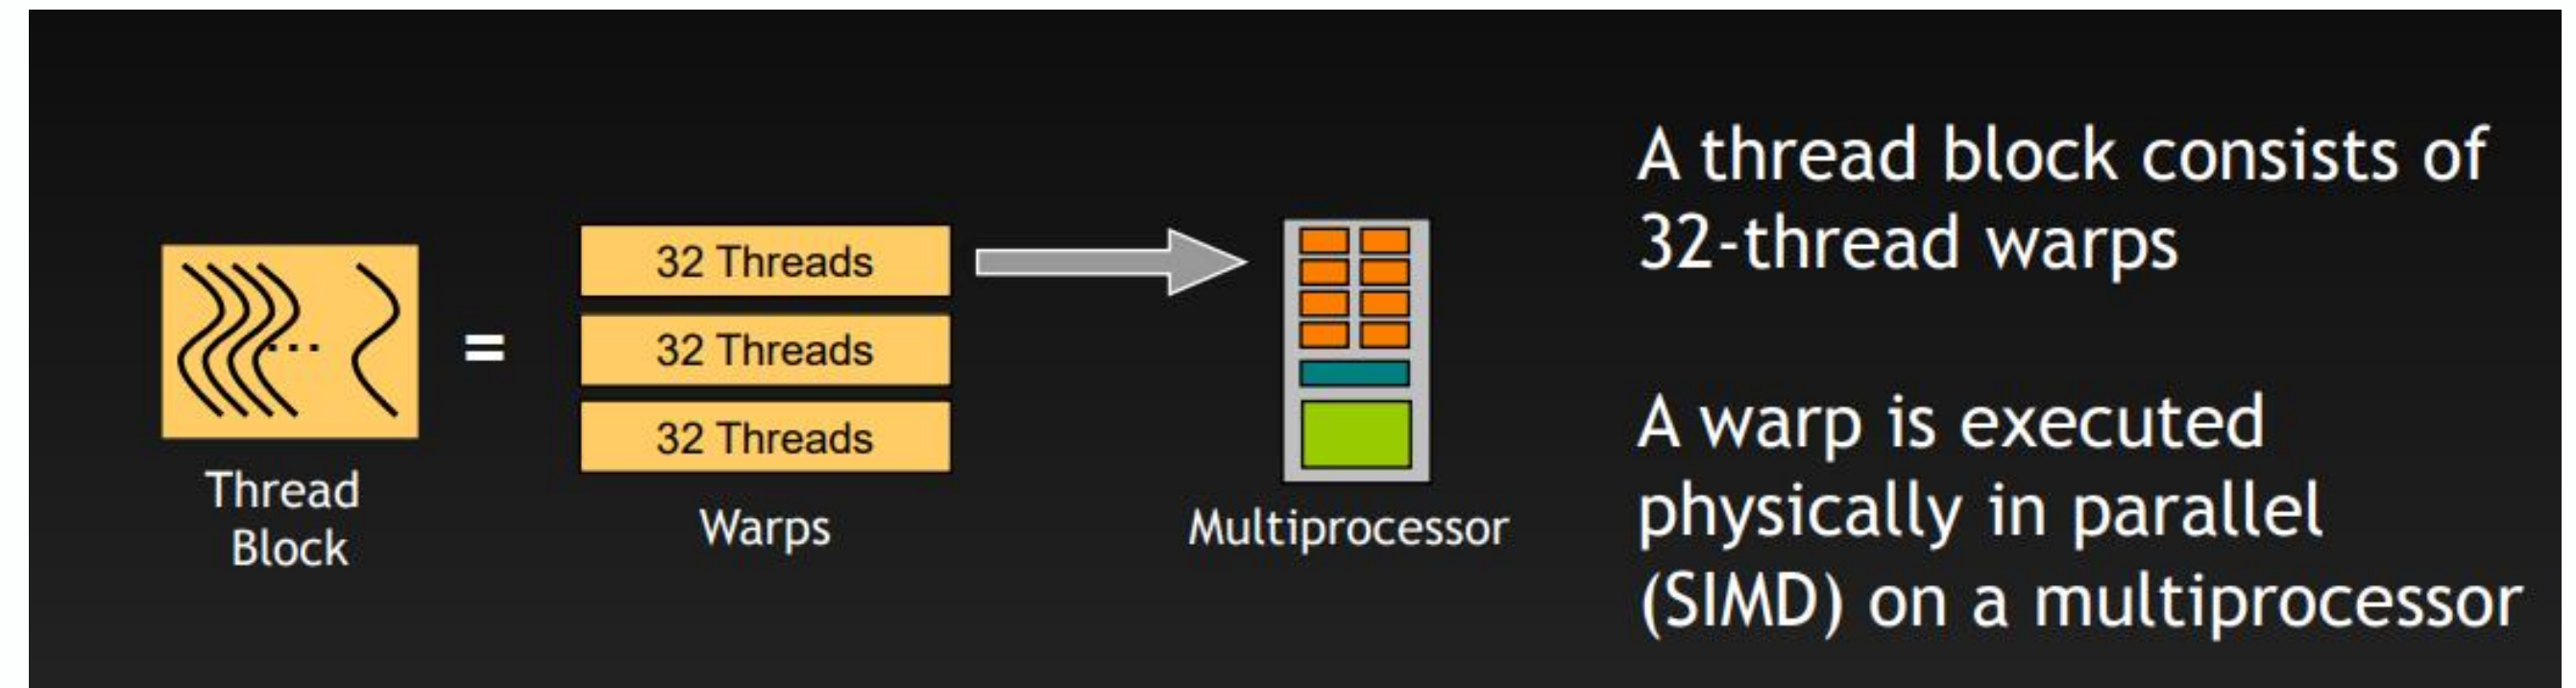
GSYNC

Results

Conclusion

Background

GPU Hardware Architecture



Summary

• Background

GSYNC


Results

Conclusion


GPU Programming Model

Introducing two new terms:

- Host



- Device



Summary

• Background

GSYNC

Results

Conclusion

Background

GPU Programming Model

```
__device__ int index(){
    return threadIdx();
}

__global__ void kernel(int n){
    printf("Hello World from %d\n",index());
    return ;
}

__host__ int main(int argc, char ** argv){

    int blocks = 2;
    int threads_per_block = 2;

    kernel<<<blocks,threads_per_block>>>();

    return 0;
}
```

And three new keywords:

- host
- global
- device

Summary

- Background

GSYNC

Results

Conclusion

Background

GPU Programming Model

```
__device__ int index(){  
    return threadIdx();  
}
```

parallel code

```
__global__ void kernel(int n){  
    printf("Hello World from %d\n",index());  
    return ;  
}
```

parallel code

```
__host__ int main(int argc, char ** argv){  
  
    int blocks = 2;  
    int threads_per_block = 2;
```

serial code

```
    kernel<<<blocks,threads_per_block>>>();
```

parallel code

```
    return 0;  
}
```

serial code

Summary

• Background

GSYNC

Results

Conclusion

GSYNC implementation:

1. Atomic functions
2. Spinlocks
3. Ticketlocks
4. Barriers
5. Merge library with M4 MACROS

Summary

Background

• GSYNC

Results

Conclusion

Atomic operations:

- FETCH_AND_ADD
- INCREMENT
- FETCH_AND_INCREMENT
- DECREMENT
- FETCH_AND_DECREMENT
- XCHG
- CMPXCHG

Summary

Background

• GSYNC

Results

Conclusion

SPINLOCK:

```
__host__ gpulock_t * gspinlock_init(gpulock_t * lock);
```

```
__device__ void gspinlock_lock(gpulock_t * lock);
```

```
__device__ void gspinlock_unlock(gpulock_t * lock);
```

Summary

Background

• GSYNC

Results

Conclusion

TICKETLOCK:

```
__host__ gputlock_t * gticketlock_init(gputlock_t * lock);
```

```
__device__ void gticketlock_lock(gputlock_t * lock);
```

```
__device__ void gticketlock_unlock(gputlock_t * lock);
```

Summary

Background

• GSYNC

Results

Conclusion


```
__host__ gbarrier_t * gbarrier_init(  
gbarrier_t * barrier,  
unsigned int init_value);
```

```
__device__ void gbarrier_wait(gbarrier_t * barrier);
```

```
__device__ void gbarrier_destroy(gbarrier_t * barrier);
```

Summary

Background

• GSYNC

Results

Conclusion

Usage Functions:

- **G_MALLOC**(destination, size)
-- malloc -> cudaMallocManaged
- **CREATE**(func,p,args...) --> **func <<< p,1 >>> (p,args)**
 - function must be first argument
 - number of total threads must be second argument
 - rest of args

Summary

Background

• GSYNC

Results

Conclusion

GSYNC & M4 MACROS

Lock Functions:

- **LOCKDEC**(lock)
- **LOCKINIT**(lock)
- **LOCK**(lock)
- **UNLOCK**(lock)

Barrier Functions:

- **BARDEC** (barrier)
- **BARINIT**(barrier, *int* num)
- **BARRIER**(barrier)

Helper Functions:

- **WAIT_FOR_END**(*int* threads)
- **GET_PID** (*int* pid)
- **GCLOCK_START**(start)
- **GCLOCK_END**(end)
- **GCLOCK_DIFF**(*float* ms_elapsed)

Summary

Background

• GSYNC

Results

Conclusion

Results

Summary
Background
GSYNC
• Results
Conclusion

My TestApplication:

- Spinlocks only:

[lock demo](#)

- Spinlocks & Barriers:

[lock&bar demo](#)

Results

```
__global__ void mykernel(
int P,
gpulock_t * lock) {
    int pid;
    GET_PID(pid);
    LOCK(lock);
    printf(" [D]  (cs%d)\n",pid);
    UNLOCK(lock);
}

__global__ void mykernel(
int P,
gpulock_t * lock,
gbarrier_t * barrier) {

    int pid;
    GET_PID(pid);
    LOCK(lock);
    printf(" [D]  (cs%d)\n",pid);
    UNLOCK(lock);
    printf(" [D]  (before %d,%u)\n",pid,*barrier);
    BARRIER(barrier);
    printf(" [D]  (after %d,%u)\n",pid,*barrier)
}
```

Results

Summary

Background

GSYNC

• Results

Conclusion

Results

****NOTE****

“ Internally printf() uses a shared data structure and so ***it is possible that calling printf() might change the order of execution of threads.*** In particular, a thread which calls printf() might take a longer execution path than one which does not call printf(), and that path length is dependent upon the parameters of the printf(). Note, however, that CUDA makes no guarantees of thread execution order except at explicit __syncthreads() barriers, so ***it is impossible to tell whether execution order has been modified by printf() or by other scheduling behaviour in the hardware.***”

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#formatted-output>

Results

Summary

Background

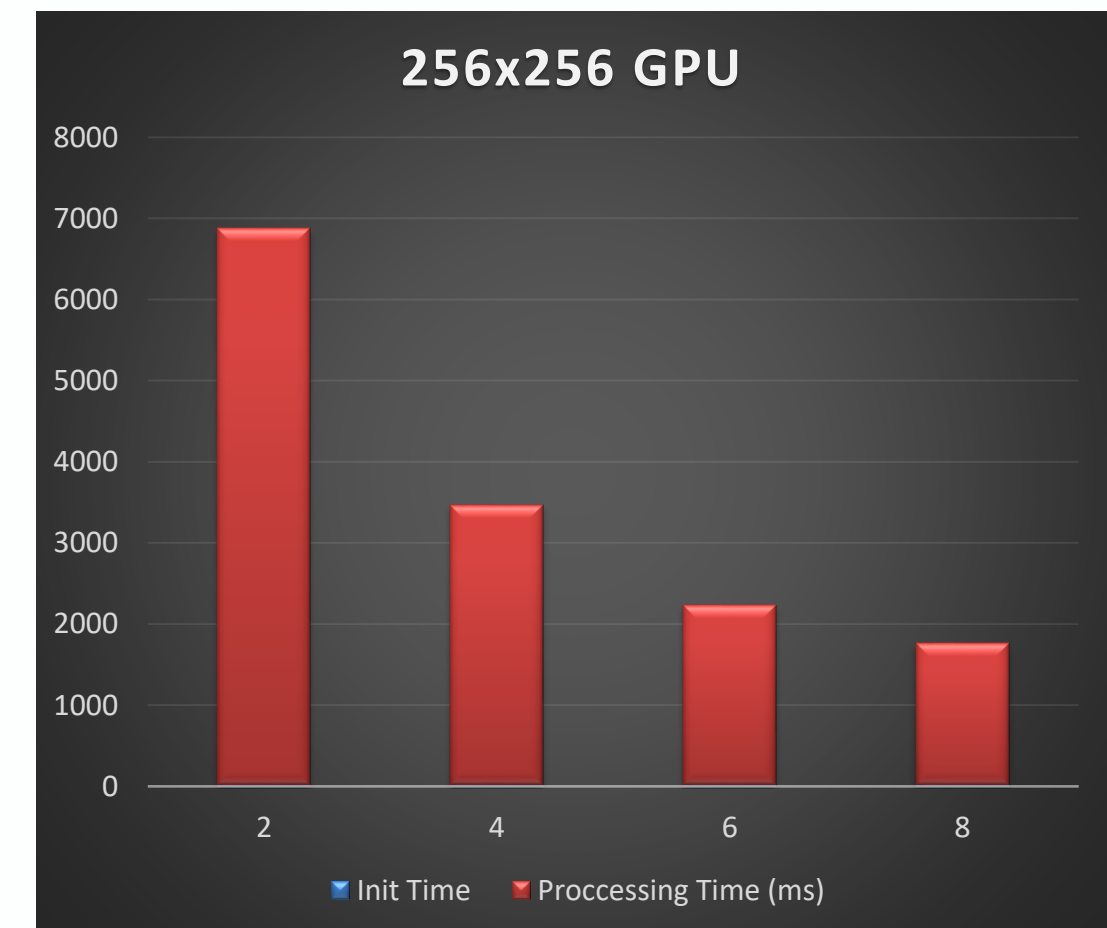
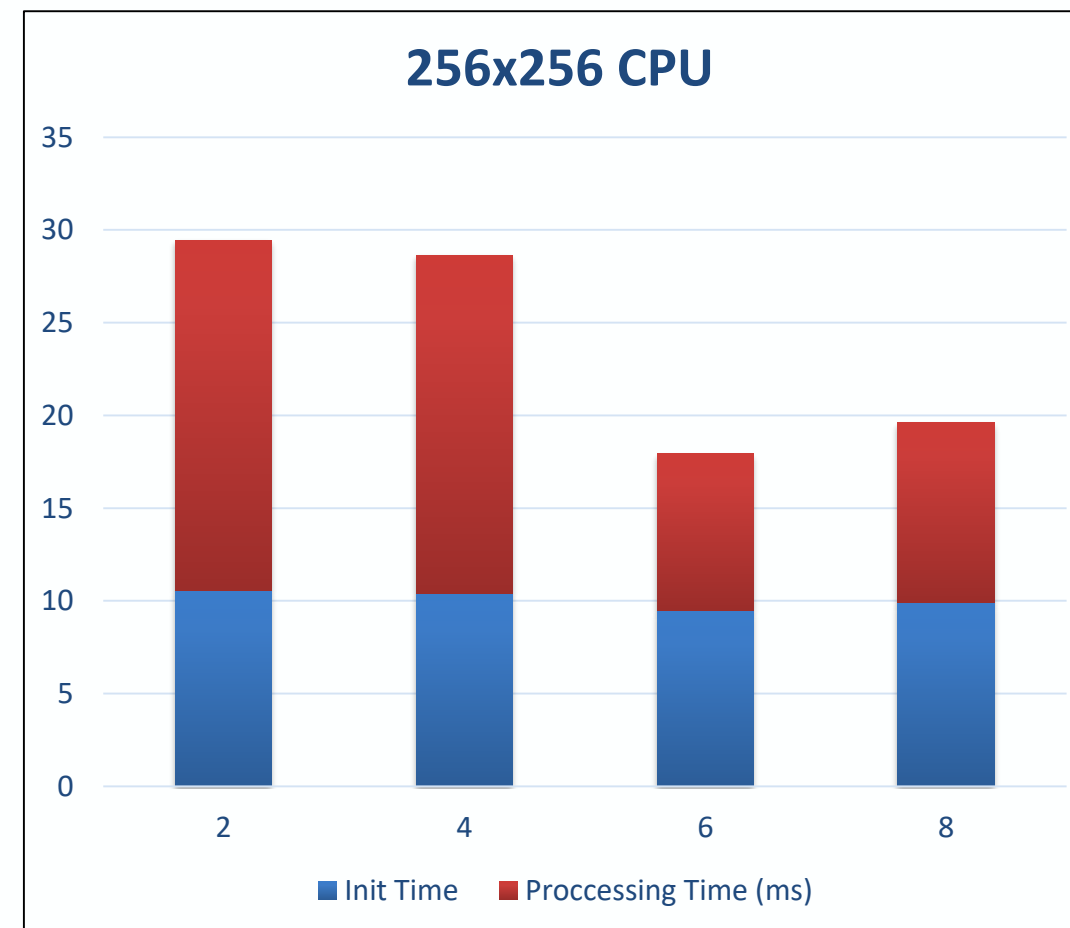
GSYNC

• Results

Conclusion

Results

GSYNC on Matrix Multiplication 256 x 256



Results

Summary

Background



GSYNC

• Results

Conclusion

Results

GSYNC on Matrix Multiplication 512 x 512




Results

- Summary
- Background
- GSYNC
- Results
- Conclusion

Results

GSYNC on Matrix Multiplication 256 x 256 (more threads)




Results

- Summary
- Background
- GSYNC
- Results
- Conclusion

Results

GSYNC on Matrix Multiplication 512 x 512 (more threads)



Results

Summary

Background

GSYNC

• Results

Conclusion

Results

Ported GSYNC **SPLASH-2** Programs:



- LU
- Radix
- FFT (working partially)

Results

- Summary
- Background
- GSYNC
- Results
- Conclusion

Results

GSYNC on LU 1024 x 1024




Results

- Summary
- Background
- GSYNC
- Results
- Conclusion

Results

GSYNC on LU 1024 x 1024 (more threads)




Results

- Summary
- Background
- GSYNC
- Results
- Conclusion

Results

GSYNC on LU 2048 x 2048 (more threads)



Conclusion

Summary

Background

GSYNC

Results

• Conclusion

Conclusion

GSYNC is a library that utilizes the atomic operations provided by the CUDA API and offers synchronization between the GPU threads.

Also emerged with the M4 Macros.

GSYNC disadvantages:

- SIMD in Blocks.
create many blocks <<< create many threads
- Access to memory allocated by *cudaMallocManaged* is slow.

Conclusion

Summary

Background

GSYNC

Results

• Conclusion

Conclusion

Do not buy GSYNC

